


# Neural Radiance Fields

# Single View 3D Reconstruction




**Objective** : Given a single view (RGB / Grayscale/ ..) , estimate the 3D geometry of the object



**Methods** : Occupancy grid estimation, Truncated Signed Distance Function (TSDF) regression, parametric models like SMAL & SMPL , etc.

# Multi-View 3D Reconstruction



**What happens when we have multiple views?**

**Objective :** Find a 3D representation which is consistent with input views.

**Application :** Rendering novel views (can observe freely from any point in 3D space)

**Methods :** Structure from Motion (SfM), COLMAP, **Neural Radiance Fields (NeRFs)**

# Citations So Far....

- Original NeRF paper - ECCV 2020 Oral - 2694 citations in 3 years

## NeRF

Representing Scenes as Neural Radiance Fields for View Synthesis

ECCV 2020 Oral - Best Paper Honorable Mention

Ben Mildenhall\*

UC Berkeley

Pratul P. Srinivasan\*

UC Berkeley

Matthew Tancik\*

UC Berkeley

Jonathan T. Barron

Google Research

Ravi Ramamoorthi

UC San Diego


Ren Ng

UC Berkeley


\*Denotes Equal Contribution

# What does a pixel Measure ?


Flux across sensor area — power per unit time



**Definition:** The field *radiance* (*luminance*) at a point in space in a given direction is the power per unit solid angle per unit area perpendicular to the direction



# What does a pixel Measure ?




$$\int \int L(\mathbf{x}, \omega) \cos \theta \, dA \, d\Omega$$

(for a very, very narrow band of  $\mathbf{x}, \omega$  — depending on sensor size and lens)

$$\propto L(\mathbf{x}^*, \omega)$$

radiance for:  $\mathbf{x}^*$  = optical centre,  $\omega$  = direction from  $\mathbf{x}^*$  to pixel sensor

# Surface-Rendering



Here, we assume that ray travels through air (without any scattering/absorption). Hence, pixel value is radiance of the surface point in that direction

$$\propto L(\mathbf{x}^*, \omega)$$

radiance for:  $\mathbf{x}^*$  = optical centre,  $\omega$  = direction from  $\mathbf{x}^*$  to pixel sensor

$$= L(\mathbf{x}^* - \lambda\omega, \omega)$$

Pixel value  $\rightarrow$  Outgoing radiance at a **single** (surface) point

# What happens when medium of transmission changes ?



Fog



Scattering




Liquids

## Exercise :

- List other scenarios where assumption that “pixel value is radiance of the surface point in a given direction”.



# Surface-Rendering



$$\propto L(\mathbf{x}^*, \omega)$$

radiance for:  $\mathbf{x}^*$  = pixel sensor centre,  $\omega$  = direction from  $\mathbf{x}$  to optical centre

$$\neq L(\mathbf{x}^* - \lambda\omega, \omega)$$

Pixel value **cannot** be reduced to emitted radiance by a single surface point

How to model  $L(\mathbf{x}, \omega)$  vary along a ray ?

# Transmittance

$$T(\mathbf{x}, \mathbf{y})$$


What fraction of radiance at  $\mathbf{x}$  in direction of  $\mathbf{y}$ , reaches  $\mathbf{y}$ ?  
(along a straight line under absorption-only model)

**Homogenous Medium:**

$$e^{-\sigma \|\mathbf{x} - \mathbf{y}\|}$$

**Non-Homogenous Medium:**


$$e^{-\int_{t=0}^{\|\mathbf{x} - \mathbf{y}\|} \sigma(\mathbf{x} + \omega \mathbf{t}) dt}$$



**Multiplicativity:**

$$T(\mathbf{x}, \mathbf{y}) = T(\mathbf{x}, \mathbf{z})T(\mathbf{z}, \mathbf{y})$$

# Absorption Only Volumetric Rendering



$$L(\mathbf{x}, \omega) = T(\mathbf{x}, \mathbf{x}_z)L(\mathbf{x}_z, \omega)$$


# Emission-Absorption Volumetric Rendering



$$L(\mathbf{x}, \omega) = T(\mathbf{x}, \mathbf{x}_z)L(\mathbf{x}_z, \omega) + \int_0^z T(\mathbf{x}, \mathbf{x}_t)\sigma(\mathbf{x}_t)L_e(\mathbf{x}_t, \omega)dt$$

Can we find this model analytically ? or Should we find an approximation ?

# Volume Rendering




$$L(\mathbf{x}, \omega) = \sum_{i=1}^N (\text{contribution from } i^{\text{th}} \text{ segment})$$

**Approximate with a discrete sum**

$\mathbf{x}_{t_i}$  :  $i^{\text{th}}$  sample along ray at depth  $t_i$

$\Delta t$  : distance between successive samples

# Volume Rendering




$$L(\mathbf{x}, \omega) = \sum_{i=1}^N T(\mathbf{x}, \mathbf{x}_{t_i}) (1 - e^{-\sigma_{t_i} \Delta t}) L_e(\mathbf{x}_{t_i}, \omega)$$

$\mathbf{x}_{t_i}$  :  $i^{\text{th}}$  sample along ray at depth  $t_i$

$\Delta t$  : distance between successive samples

# Volume Rendering



1. Draw uniform samples along a ray (N segments, or N+1 points)
2. Compute transmittance between camera and each sample
3. Aggregate contributions across segments to get overall radiance (color)

# Volume Rendering - Summary

Rendering a ray along medium :


- Computer per-point density
- Computer per-point emitted color given a direction

**If we can render a ray, we can synthesize the entire image.**

**Differentiable w.r.t density, emitted light !!!**



# Neural Radiance Fields



A scene is represented by NeRF such that :

- Given an input point, the network predicts density
- Given an input point and direction, the network predicts color

# Volumetric Rendering - Neural Radiance Fields

Rendering model for ray  $r(t) = o + td$ :

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$


weights                      colors

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment  $i$ :

$$\alpha_i = 1 - e^{-\sigma_i \delta t_i} \leftarrow \text{Density} * \text{Distance Between Points}$$




# Training - Neural Radiance Fields

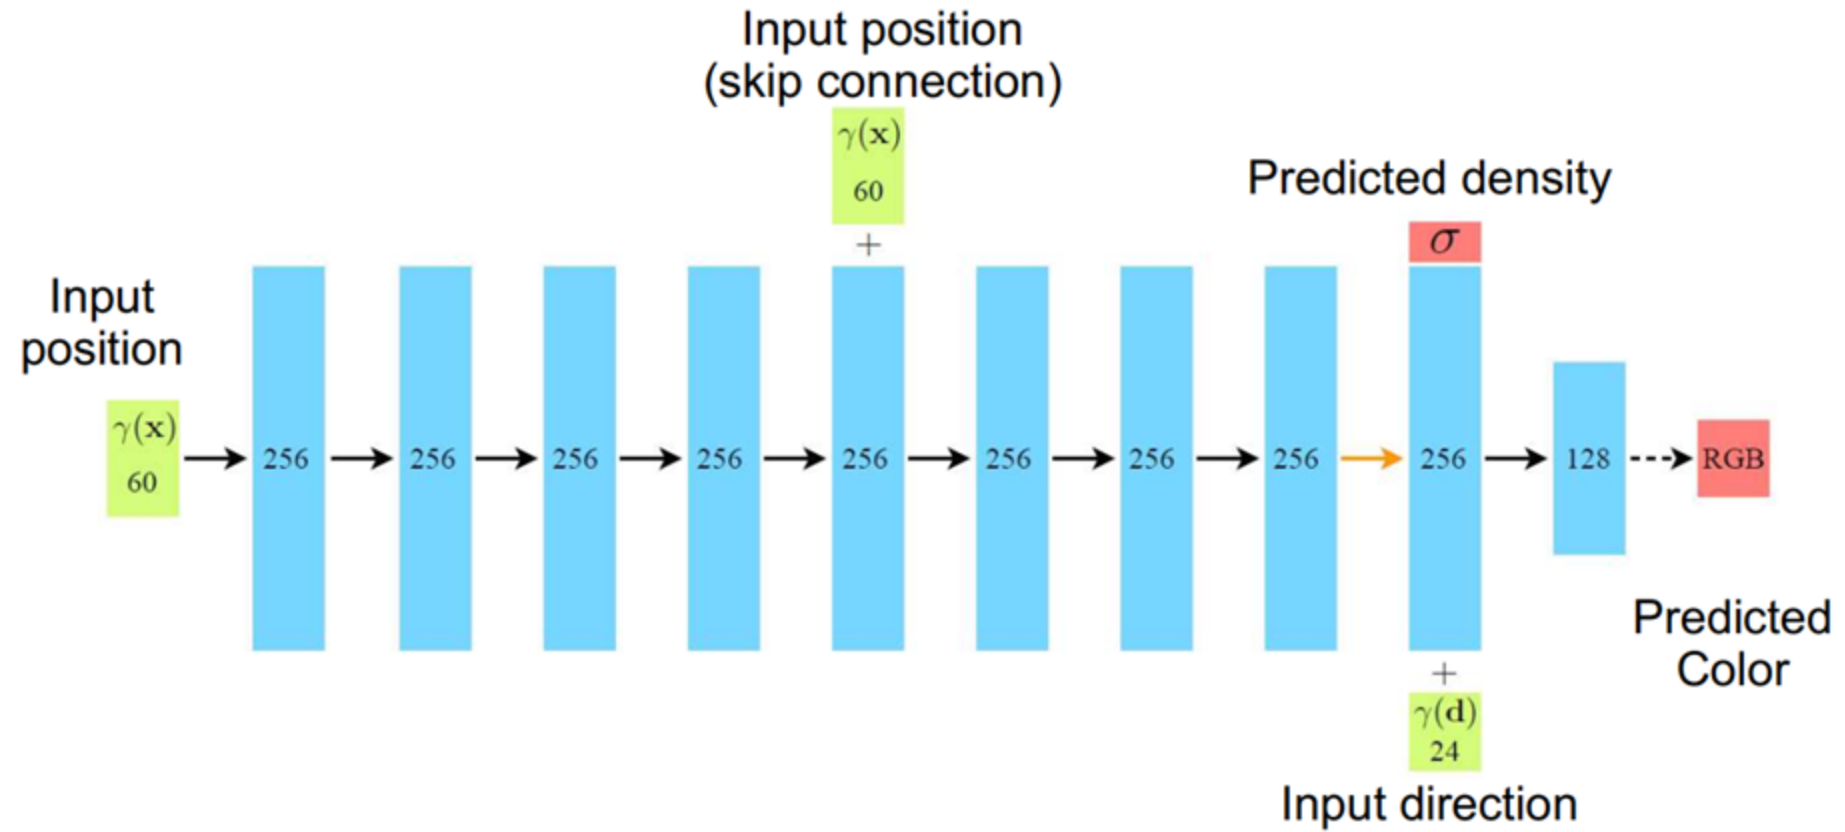
$$\min_{\theta} \sum_I \sum_{\mathbf{p}} \|\text{render}(\mathbf{p}, \pi; f_{\theta}) - I[\mathbf{p}]\|^2$$

(pixel, camera)  $\rightarrow$  ray


volume rendering



# Network Structure




# What kind of 3D representation NeRF learns ?



- It is neither a point cloud nor a 3D mesh
- It is a continuous voxel representation of the 3D scene

# Story So Far .....




- Acquire multi-view images of a scene
- Run COLMAP to extract camera-poses
- Define a NeRF :  $\mathbf{f}_{\theta}$
- Train  $\mathbf{f}_{\theta}$  with a volumetric rendering based reconstruction loss


# Initial Attempt



100 training images




Optimized Neural Net



Novel-view Renderings

# NeRF - Positional Encoding



Input Image



Using a 'standard' MLP

$\mathbf{v}$

$\sin(\mathbf{v}), \cos(\mathbf{v})$

$\sin(2\mathbf{v}), \cos(2\mathbf{v})$

$\sin(4\mathbf{v}), \cos(4\mathbf{v})$

$\dots$

$\sin(2^{L-1}\mathbf{v}), \cos(2^{L-1}\mathbf{v})$

$\gamma(\mathbf{v})$



# NeRF - Positional Encoding



Input Image






Using a 'standard' MLP



Using position encoding


# NeRF - Positional Encoding



# NeRF - Positional Encoding



100 training images




(using position encodings in input)




Novel-view Renderings

# NeRF - Uniform To Hierarchical Sampling




$$L(\mathbf{x}, \omega) = \sum_{i=1}^N T(\mathbf{x}, \mathbf{x}_{t_i}) (1 - e^{-\sigma_{t_i} \Delta t}) L_e(\mathbf{x}_{t_i}, \omega)$$

# NeRF - Uniform To Hierarchical Sampling




- Coarser Network : Sample points uniformly
- Finer Network : Sample more points depending on weights of the segments

# NeRF - Inverse Depth Sampling



- **Problem** : Need to model background for real-scenes. For e.g. sky is at infinity.
- In general, objects of interest are more closer to the camera. Hence sample, uniformly in  $(1/d)$  space; where  $d$  is depth

# Synthetic Scenes





# Real Scenes





# Depth Visualization

# Applications in Augmented Reality

# Mip-NeRF

# Mip-NeRF

- NeRF uses a single ray per-pixel and may reproduce renderings that are blurred or aliased
- Instead, we need to render multiple rays as illustrated in the purple cone Fig. (b)
- As we know rendering is expensive, it's impractical to render multiple rays for each cone
- Mip-NeRF extends NeRF to represent scenes at continuous scale by rendering anti-aliased frustums

# Mip-NeRF – Integrated Positional Encoding

Positional Encoding (PE) maps a single point into a feature vector

Integrated Positional Encoding (IPE) considers gaussians instead of infinitesimal points  
This allows “region of space” as query to a coordinated network

**When wider region is considered, contribution from higher frequencies shrinks down.**

# Mip-NeRF - Results

# Mip-NeRF - Results

# Mip-NeRF 360

## Limitations of Previous Representations :

- Both NeRF and mip-NeRF needs a bounded domain i.e they only show results on forward-facing scenes or synthetic datasets

## Problem in Representing Unbounded Scenes :

- Large scene requires more network capacity, which is expensive
- Observations are sparse and reconstruction becomes ill-posed
- How to model far away objects like distant wall, horizon ?

## Solutions:

- Apply a Kalman-like warp to mip-NeRF Gaussians
- Online distillation from large MLP to small MLP
- Regularization of density along ray intervals



# Mip-NeRF 360


Mip-NeRF casts Gaussians from camera positions. For large scenes, at far away distances we get elongated Gaussians which violates bounded assumption in mip-NeRF

Mip-NeRF 360 contracts the elongated Gaussians in the bounded space of mip-NeRF





# Dynamic Scenes – A Simple Approach



- Does-not generalize well for strictly monocular views
- Works reasonably well for multi-view dynamic scenes but fails to model disocclusions, shadows etc.

# D-NeRF: Neural Radiance Fields for Dynamic Scenes

- Maps an observed/deformed scene to canonical space using a deformable network
- Fits a radiance field for the canonical space
- Decouples motion and space by using two separate MLP networks





# Nerfies: Deformable Neural Radiance Fields

- Trace camera rays in the observed frame and transform samples along the canonical space using deformation field
- Query the template NeRF for these transformed points
- Instead of position-encoded time, they use a learnable deformation field
- Appearance code is used to handle illumination variations







# Hyper-NeRF

- Trace camera rays in the observed frame and warps these sampled points to the canonical space using deformation field
- Slice a surface from canonical hyper-space using ambient slocing network
- Concatenate them and query the template NeRF









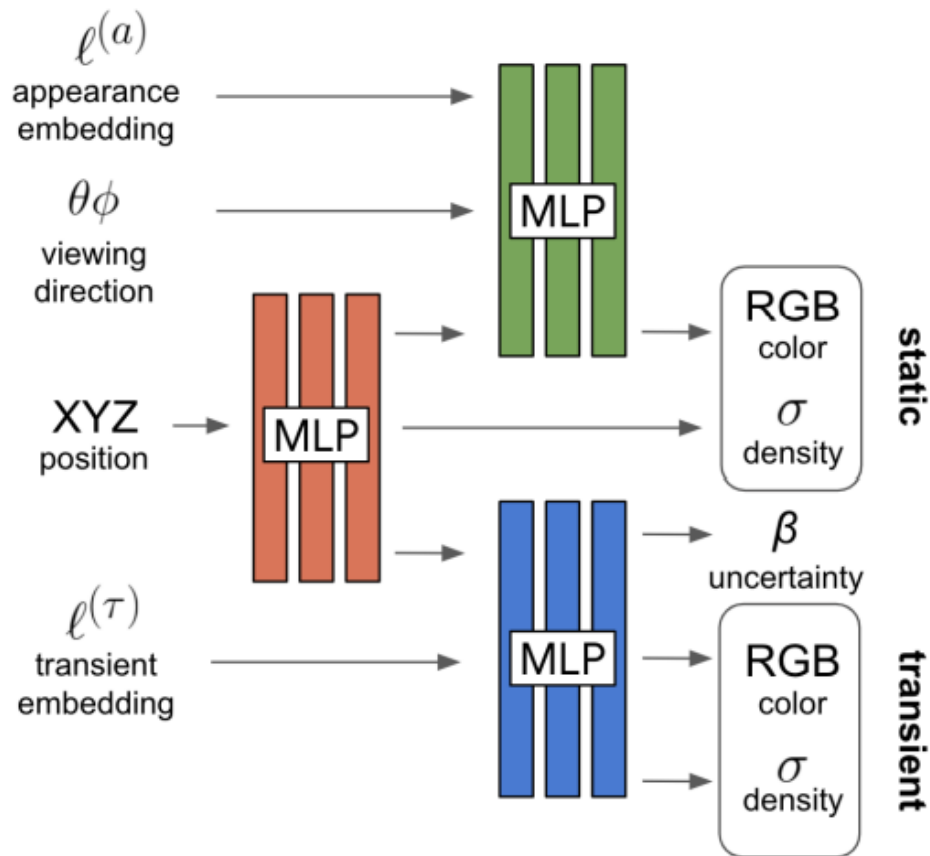




# Appearance Changes

- Exposure Differences
- Lighting Changes (Day, Night, ...)
- Passing by clouds
- External Lighting

# NeRF in the Wild



- Use learnable appearance embeddings to model exposure variations
- Use learnable transient embeddings to model transient color variations.

















# References

- [\*\*https://www.matthewtancik.com/nerf\*\*](https://www.matthewtancik.com/nerf)
- <https://learning3d.github.io/>
- <https://sites.google.com/berkeley.edu/nerf-tutorial/home>
- <https://jonbarron.info/mipnerf/>
- <https://jonbarron.info/mipnerf360/>
- <https://www.albertpumarola.com/research/D-NeRF/index.html>
- <https://nerfies.github.io/>
- <https://nerf-w.github.io/>
- <https://dreamfusion3d.github.io/>